

# DecisionTree

April 3, 2019

```
In [1]: from collections import Counter
```

```
import numpy as np
from numpy import genfromtxt
import pandas as pd
import scipy.io
from scipy import stats
import matplotlib.pyplot as plt
```

```
import random
```

```
In [2]: class DecisionTree:
```

```
    def __init__(self, max_depth=4, prune=True, threshold=.1, shuffle_features=False, n_random_features=10):
```

```
        self.left = None
        self.right = None
        self.label = None
        self.best_feature = None
        self.best_split_point = None
        self.best_gain = None
        self.n_features = None
```

```
        self.n_random_features = n_random_features
        self.shuffle_features = shuffle_features
        self.threshold = threshold
        self.prune = prune
        self.max_depth = max_depth
```

```
    @staticmethod
```

```
    def gini_impurity(y):
```

```
        counter = Counter(y)
        impurity = 1 - sum([(counter[k] / sum(counter.values()))**2 for k in counter.keys()])
```

```

        return impurity

def _split_point(self, X, y):

    best_feature = None
    best_split_point = None
    best_gain = 0

    self.n_features = X.shape[1]

    features = np.arange(self.n_features)

    # for random forest
    if self.shuffle_features:

        if self.n_random_features is None:

            self.n_random_features = int(np.sqrt(self.n_features))

        features = np.random.choice(features, self.n_random_features, replace=False)

    for col in features:

        for split_point in np.unique(X[:, col]):

            current_impurity = self.gini_impurity(y)

            left_y = y[X[:, col] < split_point]
            right_y = y[X[:, col] >= split_point]

            left_gini = self.gini_impurity(left_y)
            right_gini = self.gini_impurity(right_y)

            left_weight = left_y.shape[0] / X.shape[0]
            right_weight = right_y.shape[0] / X.shape[0]

            impurity = left_weight * left_gini + right_gini * right_weight

            # Compare current impurity with child nodes.
            # Higher gain => Better
            gain = current_impurity - impurity

            if gain > best_gain:
                best_gain = gain
                best_feature = col
                best_split_point = split_point

```

```

self.best_gain = best_gain
self.best_feature = best_feature
self.best_split_point = best_split_point

def fit(self, X, y):

    self._fit(X, y)

    if self.prune:
        self.prune_tree(y)

def _fit(self, X, y):

    if self.max_depth > 0:

        self._split_point(X, y)

        # No split is done -> splitting doesn't produce higher accuracy
        if self.best_feature is None:

            self.label = self._get_label(y)
            return

        left_X = X[X[:, self.best_feature] < self.best_split_point, :]
        left_y = y[X[:, self.best_feature] < self.best_split_point]

        right_X = X[X[:, self.best_feature] >= self.best_split_point, :]
        right_y = y[X[:, self.best_feature] >= self.best_split_point]

        # if one node has all or no points, no split is needed
        if len(left_y) == 0 or len(right_y) == 0:

            self.label = self._get_label(y)
            self.best_feature = None
            self.best_split_point = None

        else:

            self.left = DecisionTree(max_depth=self.max_depth-1, prune=self.prune,
                                     n_random_features=self.n_random_features,

            self.right = DecisionTree(max_depth=self.max_depth-1, prune=self.prune,

```

```

n_random_features=self.n_random_features

self.left._fit(left_X, left_y)
self.right._fit(right_X, right_y)

# current node reached the maximum depth
else:

    self.label = self._get_label(y)

def _get_label(self, y):
    counter = Counter(y)
    return int(max(counter.items(), key=lambda x: x[1])[0])

def predict(self, X):

    return np.array([self._predict(x) for x in X])

def _predict(self, x):

    # leaf nodes
    if self.max_depth == 0 or self.best_feature is None:

        return self.label

    else:

        if self.left is None or self.right is None:

            return self.label

        if x[self.best_feature] < self.best_split_point:

            return self.left._predict(x)

        return self.right._predict(x)

def cross_entropy(self, y, pred):

    p = sum((pred==1)) / len(pred)

```

```

        return -(y*np.log(p) + (1-y)*np.log(1-p))

def prune_tree(self, y):

    # base case : leaf node
    if self.best_feature is None:
        return

    if self.left is not None:
        self.left.prune_tree(y)

    if self.right is not None:
        self.right.prune_tree(y)

    if self.best_gain < self.threshold:
        self.left = self.right = None
        self.best_feature = self.best_split_point = self.impurity = None
        self.label = self._get_label(y)

def loss(self, y, pred):

    return self.cross_entropy(y, pred).sum() / len(y)

def accuracy(self, X, y):

    pred = self.predict(X)
    return sum(y==pred) / len(y)

def __repr__(self):

    return self._show_tree(0)

def _show_tree(self, d):

    if self.max_depth >= 0:

        indent = '  ' * d

        tree = f'\n{indent}Depth : {d}, Split Feature : {self.best_feature}, Split
        left = right = label = ''

```

```

        if self.left is not None:
            left = self.left._show_tree(d+1)

        if self.right is not None:
            right = self.right._show_tree(d+1)

        if self.left is None and self.right is None:
            return f'\n{indent}Depth : {d}, Label : {self.label}'

        else:
            return tree + left + right

class RandomForest():

    def __init__(self, max_depth=4, n_estimators=5, threshold=.02, prune=False, n_random_features=None, bootstrap=False):

        self.trees = []

        self.n_estimators = n_estimators
        self.threshold = threshold
        self.prune = prune
        self.max_depth = max_depth
        self.n_random_features = n_random_features
        self.bootstrap = bootstrap

    def fit(self, X, y):

        if self.n_random_features is None:
            self.n_random_features = int(np.sqrt(X.shape[1]))

        for i in range(self.n_estimators):

            tree = DecisionTree(max_depth=self.max_depth, prune=self.prune, threshold=self.threshold,
                                shuffle_features=True, n_random_features=self.n_random_features)

            if self.bootstrap:
                idx = np.random.choice(np.arange(X.shape[0]), int(X.shape[0]*.8))
                data, label = X[idx], y[idx]

                tree.fit(data, label)
            else:
                tree.fit(X, y)
            self.trees.append(tree)

```

```

def predict(self, X):

    pred = []

    for tree in self.trees:
        pred.append(tree.predict(X))

    pred = np.mean(pred, axis=0)

    # if higher than 0.5, it means there were more label 1 than 0.
    pred[pred >= 0.5] = 1
    pred[pred < 0.5] = 0

    return pred

def accuracy(self, X, y):

    pred = self.predict(X)

    return sum(pred == y) / len(y)

```

In [3]: `def split(X, size):`

```

    if type(size) == float:
        size = round(len(X) * size)

    dat = X.copy()

    # for reproducibility
    np.random.seed(24)

    # shuffle copied data
    np.random.shuffle(dat)

    # training_data, validation_data
    return dat[size:], dat[:size]

```

```

def fillna(cols):

```

```

    for col in cols:

        data[col] = round(data[[col]].fillna(int(round(np.mean(data[col])))))

```

In [4]: `dataset = "titanic"`

```

if dataset == "titanic":
    # Load titanic data
    path_train = 'datasets/titanic/titanic_training.csv'

```

```

data = pd.read_csv(path_train)
path_test = 'datasets/titanic/titanic_testing_data.csv'
test_data = pd.read_csv(path_test)
y = data['survived']
class_names = ["Died", "Survived"]

elif dataset == "spam":
    features = [
        "pain", "private", "bank", "money", "drug", "spam", "prescription",
        "creative", "height", "featured", "differ", "width", "other",
        "energy", "business", "message", "volumes", "revision", "path",
        "meter", "memo", "planning", "pleased", "record", "out",
        "semicolon", "dollar", "sharp", "exclamation", "parenthesis",
        "square_bracket", "ampersand"
    ]
    assert len(features) == 32

    # Load spam data
    path_train = 'datasets/spam-dataset/spam_data.mat'
    data = scipy.io.loadmat(path_train)
    X = data['training_data']
    y = np.squeeze(data['training_labels'])
    Z = data['test_data']
    class_names = ["Ham", "Spam"]

else:
    raise NotImplementedError("Dataset %s not handled" % dataset)

```

## 0.0.1 Titanic

```
In [5]: fillna(['age', 'sibsp', 'parch'])
```

```

# all col values are NaN
data = data.drop(index=705).drop(['cabin', 'ticket'], axis=1)

# index=38. Missing fare with pclass=3 and embarked at S
data.loc[data['fare'].isna(), 'fare'] = data[(data['embarked'] == 'S') & (data['pclass'] == 3)]['fare'].mean()

fare_by_embarked = data[data['pclass']==1][['fare', 'embarked']].groupby('embarked').mean()
data.loc[data['embarked'].isna(), 'embarked'] = 'S'

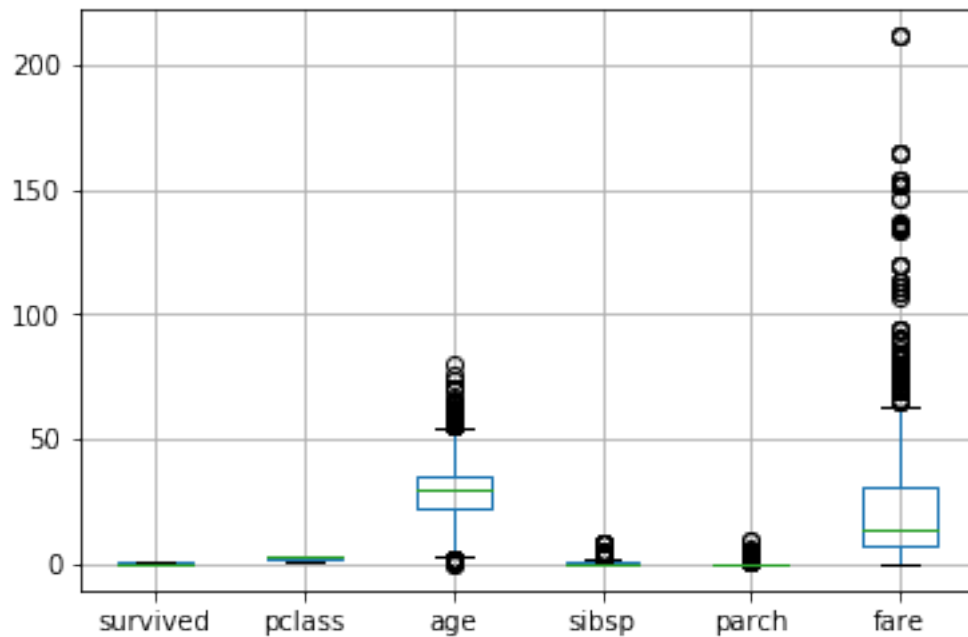
# remove outliers in fare
data = data[data['fare'] < np.percentile(data['fare'], 97.5)]

# test data
test = test_data.drop(['cabin', 'ticket'], axis=1)
test['age'] = test[['age']].fillna(data['age'].mean())

```



```
In [6]: data.boxplot();
```



```
In [7]: dat = np.array(pd.get_dummies(data.copy()))  
train, valid = split(dat, 100)
```

```
X_train_titanic, y_train_titanic = train[:, 1:], train[:, 0]  
X_val_titanic, y_val_titanic = valid[:, 1:], valid[:, 0]
```

```
In [8]: tree_titanic = DecisionTree(max_depth=2, threshold=.015)  
tree_titanic.fit(X_train_titanic, y_train_titanic)
```

```
In [9]: rf_titanic = RandomForest(max_depth=1000, n_estimators=100)  
rf_titanic.fit(X_train_titanic, y_train_titanic)
```

```
In [10]: X_test_titanic = np.array(pd.get_dummies(test))
```

```
pred = rf_titanic.predict(X_test_titanic)  
pred = pred.astype(int)
```

```
sub = pd.DataFrame(pred, columns=['Category'], index=np.arange(1, len(pred)+1, 1), dtype=int)  
sub.index.name = 'Id'
```

```
sub.to_csv('./submission.csv')
```

## 0.0.2 Spam-Ham

```
In [11]: dataset = "spam"
```


```
if dataset == "titanic":
    # Load titanic data
    path_train = 'datasets/titanic/titanic_training.csv'
    data = pd.read_csv(path_train)
    path_test = 'datasets/titanic/titanic_testing_data.csv'
    test_data = pd.read_csv(path_test)
    y = data['survived']
    class_names = ["Died", "Survived"]

elif dataset == "spam":
    features = [
        "pain", "private", "bank", "money", "drug", "spam", "prescription",
        "creative", "height", "featured", "differ", "width", "other",
        "energy", "business", "message", "volumes", "revision", "path",
        "meter", "memo", "planning", "pleased", "record", "out",
        "semicolon", "dollar", "sharp", "exclamation", "parenthesis",
        "square_bracket", "ampersand"
    ]
    assert len(features) == 32

    # Load spam data
    path_train = 'datasets/spam-dataset/spam_data.mat'
    data = scipy.io.loadmat(path_train)
    X = data['training_data']
    y = np.squeeze(data['training_labels'])
    Z = data['test_data']
    class_names = ["Ham", "Spam"]

else:
    raise NotImplementedError("Dataset %s not handled" % dataset)
```

```
In [12]: plt.boxplot(X);
```



```
In [13]: # outlier values for all features
idx = np.all(X <= np.percentile(X, 97.5, axis=0), axis=1)

X = X[idx]
y = y[idx]

In [14]: data = np.hstack((X, y.reshape(-1, 1)))
train, valid = split(data, .2)

X_train_spam, y_train_spam = train[:, :-2], train[:, -1]
X_valid_spam, y_valid_spam = valid[:, :-2], valid[:, -1]
X_test_spam = Z

In [47]: tree_spam = DecisionTree(max_depth=4, threshold=.02)
tree_spam.fit(X_train_spam, y_train_spam)

In [16]: rf_spam = RandomForest(max_depth=1000, n_estimators=300, n_random_features=10)
rf_spam.fit(X_train_spam, y_train_spam)

In [50]: pred = rf_spam.predict(X_test_spam)
pred = pred.astype(int)

sub = pd.DataFrame(pred, columns=['Category'], index=np.arange(1, len(pred)+1, 1), dtype=int)
sub.index.name = 'Id'

sub.to_csv('./submission.csv')
```

## 0.1 2.3

1. I used the mean value for numerical missing values. I removed cabin and ticket as, in my opinion, would not affect much to the outcome. After that there was no nan value in categorical features.
2. If improvement (current node's gini impurity - child nodes') was greater than .015, I stopped splitting or else kept splitting until fully grown and prune with same logic.
3. Simply used a list of decision trees each with different set of features.
4. I don't think so.
5. No... I'm out of ideas.. What a foooooool....

## 0.2 2.4 Performances

```
In [18]: print(f'Decision Tree on Titanic Training : {tree_titanic.accuracy(X_train_titanic, y_train_titanic)}')
         print(f'Decision Tree on Titanic Validation : {tree_titanic.accuracy(X_val_titanic, y_val_titanic)}')
```

```
Decision Tree on Titanic Training : 0.7938144329896907
```

```
Decision Tree on Titanic Validation : 0.75
```

```
In [19]: print(f'Random Forest on Titanic Training : {rf_titanic.accuracy(X_train_titanic, y_train_titanic)}')
         print(f'Random Forest on Titanic Validation : {rf_titanic.accuracy(X_val_titanic, y_val_titanic)}')
```

```
Random Forest on Titanic Training : 0.8751431844215349
```

```
Random Forest on Titanic Validation : 0.76
```

```
In [20]: print(f'Decision Tree on Spam Training : {tree_spam.accuracy(X_train_spam, y_train_spam)}')
         print(f'Decision Tree on Spam Validation : {tree_spam.accuracy(X_valid_spam, y_valid_spam)}')
```





```
Decision Tree on Spam Training : 0.8002658690594882
```

```
Decision Tree on Spam Validation : 0.8085106382978723
```







```
In [51]: print(f'Random Forest on Spam Training : {rf_spam.accuracy(X_train_spam, y_train_spam)}')
         print(f'Random Forest on Spam Validation : {rf_spam.accuracy(X_valid_spam, y_valid_spam)}')
```

```
Random Forest on Spam Training : 0.8348288467929544
```

```
Random Forest on Spam Validation : 0.8058510638297872
```

26	Alex Kassil		0.86021	4	2d
27	Han Song		0.86021	3	2d
<b>Your Best Entry</b>  Your submission scored 0.86021, which is not an improvement of your best score. Keep trying!					
28	Avery		0.86021	2	1d

titanic.png

212	tttzzz		0.78429	4	15h
213	Jacky Chow		0.78258	6	3d
214	Han Song		0.78258	5	2m
<b>Your Best Entry</b>  Your submission scored 0.75697, which is not an improvement of your best score. Keep trying!					
215	Jeremy Chui		0.78258	1	1d
216	Sathvik Nair		0.78258	1	9h

spam

### 0.3 2.5

```
In [22]: single_spam = X_train_spam[30]
single_spam
```

```
Out[22]: array([0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 1.,
                0., 0., 1., 0., 0., 0., 0., 0., 0., 3., 0., 0., 1., 0.]
```

```
In [23]: dec = DecisionTree(max_depth=4, threshold=.01)
dec.fit(X_train_spam, y_train_spam)
```

```
In [24]: dec
```

```
Out[24]:
Depth : 0, Split Feature : 28, Split Point : 1.0, Gain : 0.053420098256930226
  Depth : 1, Split Feature : 29, Split Point : 1.0, Gain : 0.013287901979625094
    Depth : 2, Split Feature : 19, Split Point : 1.0, Gain : 0.019579639101447888
      Depth : 3, Label : 0
      Depth : 3, Label : 0
      Depth : 2, Label : 0
    Depth : 1, Split Feature : 19, Split Point : 1.0, Gain : 0.046734117230732675
      Depth : 2, Split Feature : 29, Split Point : 1.0, Gain : 0.020001679772411518
        Depth : 3, Split Feature : 16, Split Point : 1.0, Gain : 0.021617585724248656
          Depth : 4, Label : 1
          Depth : 4, Label : 0
        Depth : 3, Split Feature : 3, Split Point : 1.0, Gain : 0.0404505000061805
          Depth : 4, Label : 0
          Depth : 4, Label : 1
        Depth : 2, Label : 0
```

```
In [25]: np.where(single_spam > 0)
```

```

Out [25]: (array([16, 19, 26, 29], dtype=int64),)

In [26]: np.array(features)[np.where(single_spam>0)]

Out [26]: array(['volumes', 'meter', 'dollar', 'parenthesis'], dtype='<U14')

In [27]: features[28]

Out [27]: 'exclamation'

In [28]: dec.predict(single_spam.reshape(1, -1))

Out [28]: array([0])

1. exclamation <= 1 (moved to left node)
2. parenthesis <= 1 (moved to left)
3. parenthesis <= 1 (moved to left)
4. It is not a spam

In [29]: y_train_spam[30]

Out [29]: 0.0

```

## 0.4 2.6

```

In [30]: titanic_train, titanic_valid = split(dat, .2)

In [31]: titanic_train.shape, titanic_valid.shape

Out [31]: ((778, 11), (195, 11))

In [32]: depth = np.arange(1, 50, 5)
         acc = []


         for d in depth:

             dec = DecisionTree(max_depth=d, threshold=.015, prune=False)
             dec.fit(titanic_train[:, :-2], titanic_train[:, -1])

             acc.append(dec.accuracy(titanic_valid[:, :-2], titanic_valid[:, -1]))

In [33]: plt.figure(figsize=(8,6))
         plt.plot(depth, acc)
         plt.xlabel('depth', size=15)
         plt.ylabel('Validation Accuracy', size=15);

```



Without pruning, it peaks around before the depth of 8~10 and after that, it decreases and again after some depth, it has a constant accuracy score. This is due to overfitting and after some number of depths, it will not increase the accuracy.

In [ ]: ~