

## 1 Connections between OLS, Ridge Regression, TLS, PCA, and CCA

We will review several topics we have learned so far: ordinary least-squares, ridge regression, total least squares, principle component analysis, and canonical correlation analysis. We emphasize their basic attributes, including the objective functions and the explicit form of their solutions. We will also discuss the connections and distinctions between these methods.

- What are the objective functions and closed-form solutions to OLS, ridge regression, and TLS? How do the probabilistic interpretations vary?
- Consider the matrix inversion in the solution to OLS, ridge regression, and TLS. How do the eigenvalues compare to those of the matrix  $\mathbf{X}^\top \mathbf{X}$ ?
- Suppose you have a data matrix  $\mathbf{X} \in \mathbb{R}^{n \times d}$  and output  $\mathbf{y} \in \mathbb{R}^{n \times 1}$ . Use PCA to compute the first  $k$  principal components of  $[\mathbf{X} \ \mathbf{y}]$ . Describe how this solution would relate to performing TLS on the problem.
- Suppose you have a problem of fitting  $\mathbf{y}$  from  $\mathbf{X} \in \mathbb{R}^{n \times d}$ .  $d$  is large in this case, and you want to reduce the dimensionality before doing the fitting. You have two choices for the dimensionality reduction: the first one being do PCA on  $\mathbf{X}$  and the second one being use CCA between  $\mathbf{X}$  and  $\mathbf{y}$  to reduce dimension of  $\mathbf{X}$  with  $\mathbf{X}\mathbf{U}_p$ , where  $\mathbf{U}_p \in \mathbb{R}^{d \times p}$  is the  $p$ -dimensional projection matrix solved using CCA. Which of them is better in general?

## 2 Linear regression (16 points)

In this problem, we study the loss function for ridge regression:

$$\frac{1}{2} \|\mathbf{X}\mathbf{w} - \mathbf{y}\|_2^2 + \frac{\lambda}{2} \|\mathbf{w}\|_2^2, \quad (1)$$

where  $\mathbf{X}$  a data matrix in  $\mathbb{R}^{n \times d}$  and  $\mathbf{y}$  is the response in  $\mathbb{R}^n$ . Let the regularization weight  $\lambda > 0$ . Recall that the closed-form solution to ridge regression is

$$\hat{\mathbf{w}} = (\mathbf{X}^\top \mathbf{X} + \lambda \mathbf{I}_d)^{-1} \mathbf{X}^\top \mathbf{y}. \quad (2)$$

where  $\mathbf{I}_d \in \mathbb{R}^{d \times d}$  is an identity matrix of dimension  $d$ .

- (8 points) Augment the matrix  $\mathbf{X}$  with  $d$  additional rows  $c\mathbf{e}_1^\top, c\mathbf{e}_2^\top, \dots, c\mathbf{e}_d^\top$  to get the matrix  $\mathbf{X}' \in \mathbb{R}^{n+d \times d}$ , where  $c$  is a given constant and  $\mathbf{e}_i^\top$  is a unit vector whose  $i$ th element is 1 and the rest of the

elements are zero, and augment  $y$  with  $d$  zeros to get  $y' \in \mathbb{R}^{n+d}$ :

$$X' = \begin{bmatrix} X \\ ce_1^T \\ ce_2^T \\ \vdots \\ ce_d^T \end{bmatrix} \quad y' = \begin{bmatrix} y \\ 0 \\ 0 \\ \vdots \\ 0 \end{bmatrix}$$

**Write out the closed form solution for  $w'$  for the ordinary least squares problem on  $X', y'$ .**

**Derive the concrete value of  $c$  that corresponds to the ridge regression problem (1) in terms of  $\lambda$ .** Conclude that the ridge regression estimate for  $w$  can be obtained by doing ordinary least squares on an augmented dataset.

- (b) (8 points) **Prove that applying gradient descent on the ridge regression loss function with a suitable fixed step size results in geometric convergence.** If  $X^\top X$  has maximum and minimum eigenvalues  $M$  and  $m$ , **what fixed step size should we choose as a function of  $M$ ,  $m$ , and ridge weight  $\lambda$ ?**

**Hint:** Reduce the problem to ordinary least squares. Recall that applying gradient descent on the least squares problem

$$\min_x f(x) = \min_x \frac{1}{2} \|Ax - b\|_2^2$$

results in geometric convergence when  $A^\top A$  is positive definite and using a constant step size  $\gamma = \frac{2}{\lambda_{\min}(A^\top A) + \lambda_{\max}(A^\top A)}$ . Geometric convergence means that  $f(x_k) - f(x^*) \leq c' Q^k$  for some  $0 \leq Q < 1$  and for some  $c' > 0$ . You may use this result without proof.

### 3 Classification and Separability

In this problem, we will explore how adding the label as a feature affects separability of a dataset.

- (a) Consider a dataset  $\{x_i, y_i\}_{i=1}^n$ , where  $x_i \in \mathbb{R}^d$ ,  $y \in \{-1, 1\}$  which is *not* linearly separable.

Now for each point in the dataset, consider augmenting each feature vector with the label, so that we get a new feature vector  $z_i \in \mathbb{R}^{d+1}$ , where:

$$z_i = \begin{bmatrix} x_i \\ y_i \end{bmatrix}$$

Show that using the augmented feature vectors, the dataset is linearly separable.

- (b) How will we expect the linear decision boundary defined above to perform when applied to a new unseen set of features  $\mathbf{X}^{new} \in \mathbb{R}^{n_2 \times d}$ ?

## 4 Kernel PCA (24 points)

In lectures, discussion, and homework, we learned how to use PCA to do dimensionality reduction by projecting the data to a subspace that captures most of the variability. This works well for data that is roughly Gaussian shaped, but many real-world high dimensional datasets have underlying low-dimensional structure that is not well captured by linear subspaces. However, when we lift the raw data into a higher-dimensional feature space by means of a nonlinear transformation, the underlying low-dimensional structure once again can manifest as an approximate subspace. Linear dimensionality reduction can then proceed. As we have seen in class so far, kernels are an alternate way to deal with these kinds of nonlinear patterns without having to explicitly deal with the augmented feature space. This problem asks you to discover how to apply the “kernel trick” to PCA.

Let  $\mathbf{X} \in \mathbb{R}^{n \times \ell}$  be the data matrix, where  $n$  is the number of samples and  $\ell$  is the dimension of the raw data. Namely, the data matrix contains the data points  $\mathbf{x}_j \in \mathbb{R}^\ell$  as rows

$$\mathbf{X} = \begin{pmatrix} \mathbf{x}_1^\top \\ \mathbf{x}_2^\top \\ \vdots \\ \mathbf{x}_n^\top \end{pmatrix} \in \mathbb{R}^{n \times \ell}. \quad (3)$$

- (a) (5 pts) **Compute  $\mathbf{X}\mathbf{X}^\top$  in terms of the singular value decomposition  $\mathbf{X} = \mathbf{U}\mathbf{\Sigma}\mathbf{V}^\top$  where  $\mathbf{U} \in \mathbb{R}^{n \times n}$ ,  $\mathbf{\Sigma} \in \mathbb{R}^{n \times \ell}$  and  $\mathbf{V} \in \mathbb{R}^{\ell \times \ell}$ .** Notice that  $\mathbf{X}\mathbf{X}^\top$  is the matrix of pairwise Euclidean inner products for the data points. **How would you get  $\mathbf{U}$  if you only had access to  $\mathbf{X}\mathbf{X}^\top$ ?** You may assume that the data is centered: That is assume  $\frac{1}{N} \sum_{i=1}^N x_i = 0$
- (b) (7 pts) Given a new test point  $\mathbf{x}_{test} \in \mathbb{R}^\ell$ , one central use of PCA is to compute the projection of  $\mathbf{x}_{test}$  onto the subspace spanned by the  $k$  top singular vectors  $\mathbf{v}_1, \mathbf{v}_2, \dots, \mathbf{v}_k$ .

**Express the scalar projection  $z_j = \mathbf{v}_j^\top \mathbf{x}_{test}$  onto the  $j$ -th principal component as a function of the inner products**

$$\mathbf{X}\mathbf{x}_{test} = \begin{pmatrix} \langle \mathbf{x}_1, \mathbf{x}_{test} \rangle \\ \vdots \\ \langle \mathbf{x}_n, \mathbf{x}_{test} \rangle \end{pmatrix}. \quad (4)$$

Assume that all diagonal entries of  $\mathbf{\Sigma}$  are nonzero and non-increasing, that is  $\sigma_1 \geq \sigma_2 \geq \dots > 0$ .

*Hint: Express  $\mathbf{V}^\top$  in terms of the singular values  $\mathbf{\Sigma}$ , the left singular vectors  $\mathbf{U}$  and the data matrix  $\mathbf{X}$ . If you want to use the compact form of the SVD, feel free to do so.*

- (c) (12 pts) How would you define kernelized PCA for a general kernel function  $k(\mathbf{x}_i, \mathbf{x}_j)$  (to replace the Euclidean inner product  $\langle \mathbf{x}_i, \mathbf{x}_j \rangle$ )? You may assume that the data remains centered after the feature map  $\phi$  induced by the kernel function  $k$ . That is assume  $\frac{1}{N} \sum_{i=1}^N \phi(x_i) = 0$

**Describe this in terms of a procedure which takes as inputs the training data points  $\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_n \in \mathbb{R}^\ell$  and the new test point  $\mathbf{x}_{test} \in \mathbb{R}^\ell$ , and outputs the analog of the previous part's  $z_j$  coordinate in the kernelized PCA setting. You should include how to compute  $\mathbf{U}$  from the data, as well as how to compute the analog of  $\mathbf{X}\mathbf{x}_{test}$  from the previous part.**

Invoking the SVD or computing eigenvalues/eigenvectors is fine in your procedure, as long as it is clear what matrix is having its SVD or eigenvalues/eigenvectors computed. The kernel  $k(\cdot, \cdot)$  can be used as a black-box function in your procedure as long as it is clear what arguments it is being given.

## 5 Gradient Descent and Convexity

Here is the **Smoothed version of the Hinge loss function** (with parameter  $t$ ):

$$f(y) = \begin{cases} \frac{1}{2} - ty & \text{if } ty \leq 0, \\ \frac{1}{2}(1 - ty)^2 & \text{if } 0 < ty < 1, \\ 0 & \text{if } 1 \leq ty \end{cases}$$

Given  $x_1, x_2, \dots, x_n$  data points and  $y_1, y_2, \dots, y_n$  labels. we define the optimization problems as follows:

$$\min_w \frac{1}{n} \sum_{i=1}^n f(w^\top x_i - y_i)$$

Define:  $L(w) = \frac{1}{n} \sum_{i=1}^n f(w^\top x_i - y_i)$

- Is  $L(w)$  convex?
- Write out the Gradient Descent update equation.
- Write out the Stochastic Gradient Descent update equation.

Doodle page! Draw us something if you want or give us suggestions or complaints.  
You can also use this page to report anything suspicious that you might have noticed.