

HW1

January 28, 2019

```
In [1]: import mxnet as mx
        from mxnet import nd
        from mxnet.gluon import nn
        import numpy as np

        import time
```

0.1 1.1

0.1.1 1

```
In [98]: A = nd.random.normal(0, 1, shape=(4096, 4096))
        B = nd.random.normal(0, 1, shape=(4096, 4096))
```

2

```
In [99]: tic = time.time()

        C = nd.dot(A, B)
        C.wait_to_read()

        print(time.time() - tic)
```

0.4232213497161865

3

```
In [18]: tic = time.time()

        for i in range(4096):
            if i == 0:
                C = nd.dot(A, B[:, :1])
            else:
                C = nd.concat(C, nd.dot(A, B[:, i:i+1]))

            C.wait_to_read()

        print(time.time() - tic)
```

129.4343912601471

4

```
In [19]: tic = time.time()

        for i in range(4096):
            for j in range(4096):
                if j == 0:
                    row = nd.dot(A[i:i+1, :], B[:, :1])
                else:
                    row = nd.concat(row, nd.dot(A[i:i+1, :], B[:, j:j+1]))
            if i == 0:
                C = row
            else:
                C = nd.concat(C, row, dim=0)

        C.wait_to_read()

        print(time.time() - tic)
```

1747.2818703651428

5 For a small data which can be fit into CPU without overflowing, CPU will be faster. However, if the data doesn't fit into the buffers and is big enough, using GPU will be faster as GPU can store more data and compute asynchronous.

```
In [6]: A = nd.random.normal(0, 1, shape=(4096, 4096), ctx=mx.gpu(0))
        B = nd.random.normal(0, 1, shape=(4096, 4096), ctx=mx.gpu(0))

        tic = time.time()

        for i in range(4096):
            for j in range(4096):
                if j == 0:
                    row = nd.dot(A[i:i+1, :], B[:, :1])
                else:
                    row = nd.concat(row, nd.dot(A[i:i+1, :], B[:, j:j+1]))
            if i == 0:
                C = row
            else:
                C = nd.concat(C, row, dim=0)

        C.wait_to_read()

        print(time.time() - tic)
```

1669.3835248947144

0.1.2 1.2

1 To prove $B = ADA^T$ is PSD, we need to check if all the diagonal entries are equal or greater than 0 or that $ADA^T \geq 0$. We can rewrite them in vector form as $\sum_i a_i d_i a_i^T \geq 0$. This again becomes $\sum_i a_i^2 d_i$ and since each vector of d is nonnegative and for any real value of a , $a^2 \geq 0$, the sum will be at least 0 or greater, proving that B or ADA^T is a positive semidefinite.

2 We know that the matrix B is symmetric and that ADA^T is the eigendecomposition of it. By the properties of eigenvectors of a symmetric matrix, we know that those are orthogonal to each other. Hence we prefer to use ADA^T when we would like to get the matrix to the power of any real number. For example, it is better to compute $(ADA^T)^2$ which just becomes AD^2A^T than B^2 . Since D is a diagonal matrix, the square matrix or any powers of it is easier to compute than doing B^n as it computes some redundant procedures. Other than that, it would be better to use B .

0.1.3 1.3

In [20]: !nvidia-smi

Mon Jan 28 20:05:14 2019

+-----+									
NVIDIA-SMI 410.48					Driver Version: 410.48				
+-----+									
GPU Name		Persistence-M		Bus-Id		Disp.A		Volatile Uncorr. ECC	
Fan Temp Perf		Pwr:Usage/Cap		Memory-Usage		GPU-Util		Compute M.	
+=====+									
0 GeForce GTX 1080		Off		00000000:65:00.0		On		N/A	
26% 32C P8		11W / 180W		696MiB / 8116MiB		9%		Default	
+-----+									

Processes:						GPU Memory
GPU	PID	Type	Process name			Usage
0	1243	G	/usr/lib/xorg/Xorg			18MiB
0	1296	G	/usr/bin/gnome-shell			49MiB
0	1566	G	/usr/lib/xorg/Xorg			284MiB
0	1697	G	/usr/bin/gnome-shell			140MiB
0	2064	G	...quest-channel-token=8467226019251121638			201MiB

```
In [21]: D = nd.random.normal(0, 1, shape=(2, 2), ctx=mx.gpu(0))
D
```

```
Out[21]:
[[ -1.3204551   0.68232244]
 [-0.9858383   0.01992839]]
<NDArray 2x2 @gpu(0)>
```

0.1.4 1.4

```
In [3]: A = nd.random.normal(0, 1, shape=(4096, 4096))
B = nd.random.normal(0, 1, shape=(4096, 4096))
```

Computing directly to c_i

```
In [4]: tic = time.time()

c = np.zeros(4096)

for i in range(4096):

    c[i] = (nd.dot(A, B[:, i]).norm()**2).asscalar()

print(time.time() - tic)

15.849517583847046
```

With intermediate storage of NDArray

```
In [5]: tic = time.time()

d = nd.empty(4096)

for i in range(4096):

    nd.norm(nd.dot(A, B[:, i]), out=d[i])
    d[i] = d[i]**2

d.wait_to_read()
d = d.asnumpy()
print(time.time() - tic)

13.085696935653687
```

0.1.5 1.5

1

```
In [36]: A = nd.arange(12).reshape((3,3))
        B = nd.arange(12).reshape((3,3))
        C = nd.arange(12).reshape((3,3))

        C[:] = nd.dot(A, B) + C

        C
```

Out [36]:

```
[[ 15.  19.  23.]
 [ 45.  58.  71.]
 [ 75.  97. 119.]]
<NDArray 3x3 @cpu(0)>
```

2

```
In [55]: A = nd.arange(12).reshape((3,3))
        B = nd.arange(12).reshape((3,3))
        C = nd.arange(12).reshape((3,3))

        nd.elemwise_add(nd.dot(A, B), C, out=C)
```

Out [55]:

```
[[ 15.  19.  23.]
 [ 45.  58.  71.]
 [ 75.  97. 119.]]
<NDArray 3x3 @cpu(0)>
```

0.1.6 1.6

```
In [24]: x = nd.array(np.around(np.arange(-10, 10.1, 0.1), decimals=1)).reshape((-1,1))
        y = nd.arange(1, 21).reshape((1, -1))

        A = x**y

        A
```

Out [24]:

```
[[-1.0000000e+01  1.0000000e+02 -1.0000000e+03 ...  9.9999998e+17
 -1.0000000e+19  1.0000000e+20]
 [-9.8999996e+00  9.8009995e+01 -9.7029889e+02 ...  8.3451318e+17
 -8.2616803e+18  8.1790629e+19]
 [-9.8000002e+00  9.6040001e+01 -9.4119208e+02 ...  6.9513558e+17
 -6.8123289e+18  6.6760824e+19]
```

```
...
[ 9.8000002e+00  9.6040001e+01  9.4119208e+02 ...  6.9513558e+17
  6.8123289e+18  6.6760824e+19]
[ 9.8999996e+00  9.8009995e+01  9.7029889e+02 ...  8.3451318e+17
  8.2616803e+18  8.1790629e+19]
[ 1.0000000e+01  1.0000000e+02  1.0000000e+03 ...  9.9999998e+17
  1.0000000e+19  1.0000000e+20]]
<NDArray 201x20 @cpu(0)>
```